




ECE 8823 A / CS 8803 - ICN Interconnection Networks Spring 2017



http://tusharkrishna.ece.gatech.edu/teaching/icn_s17/


Lecture 5: Deadlocks - I

Tushar Krishna

Assistant Professor
School of Electrical and Computer Engineering
Georgia Institute of Technology

tushar@ece.gatech.edu

Announcement: Lab 2



16-router NoC

Each router connected to one
CPU + one Directory

(--num-cpus=16,
--num-dirs=16)

Just like Mesh_XY.py

**DO NOT CONNECT 4 CPU/DIR to
ONE ROUTER LIKE THE PAPER**

Link Weights:

1 for X-direction links

2 for Y-direction links

Taxonomy of Routing Algorithms

■ Classification I: path length

■ **Minimal:** shortest paths

- Example: Greedy over Ring, XY over Mesh

■ **Non-minimal:** non-shortest paths

- Example: Random and Adaptive over Ring/Mesh

Taxonomy of Routing Algorithms

- **Classification II: path diversity** (how to select between the set of all possible paths R_{xy} from the source x to the dest y)
 - **Deterministic:** always choose the same route between x and y , even if $|R_{xy}| > 1$
 - **Example:** Greedy over Ring, XY over Mesh
 - + Easy to Implement
 - - Inefficient use of bandwidth
 - **Oblivious:** choose any of the routes in R_{xy} without considering any information about current network state (i.e., congestion)
 - **Example:** Random over Ring, O1Turn over Mesh
 - + More path diversity
 - - Can lead to deadlocks (this lecture)
 - **Adaptive:** choose one of the routes in R_{xy} depending on the current network state (i.e., congestion)
 - **Example:** Adaptive over Ring/Mesh
 - + Best use of available bandwidth
 - - Need to track congestion, can lead to deadlocks

Taxonomy of Routing Algorithms

- **Classification III – implementation**

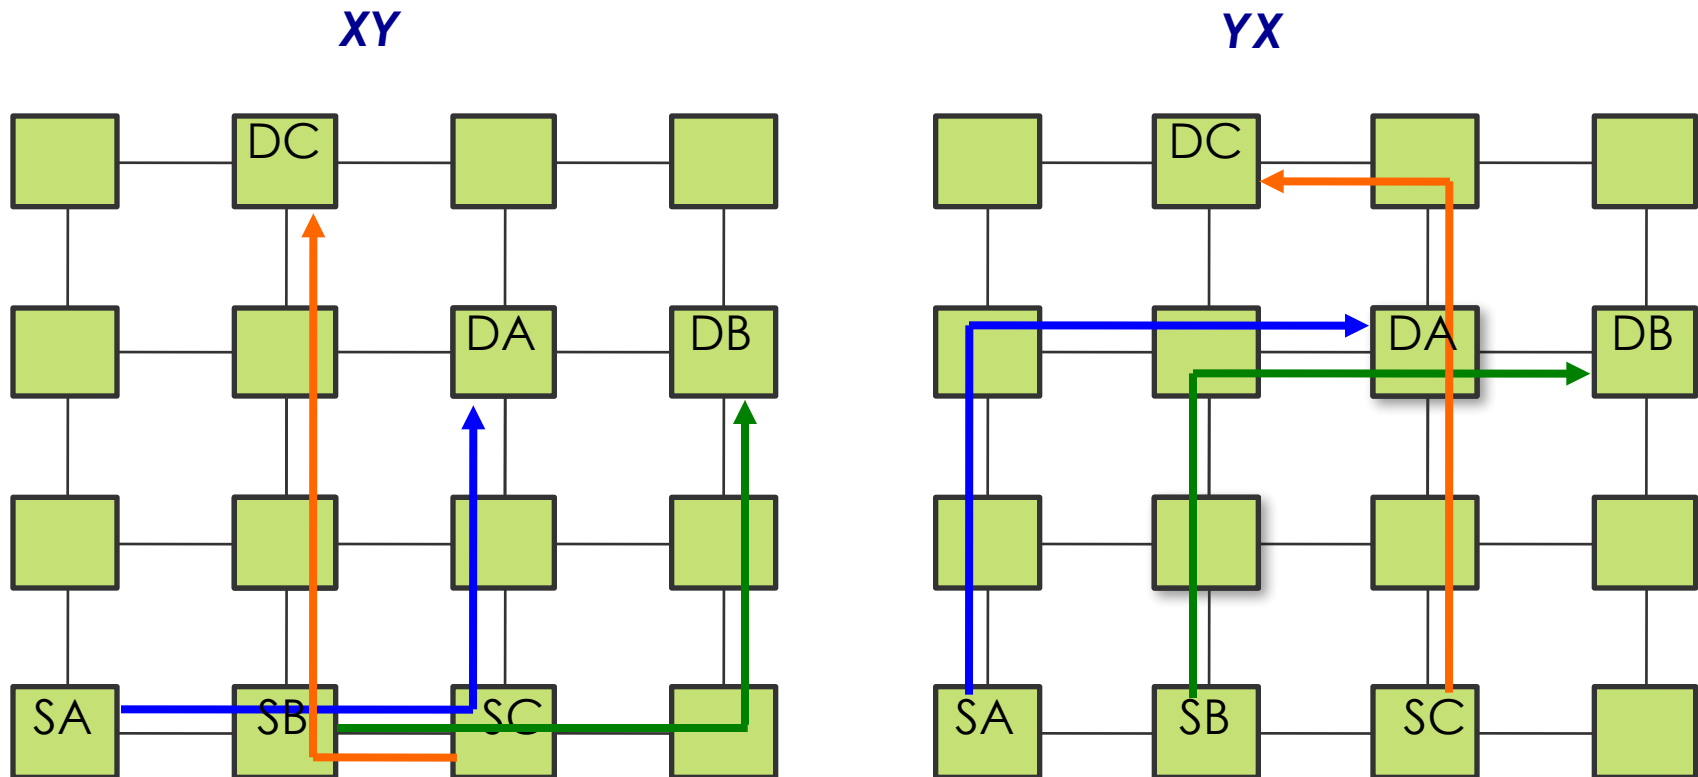
- **Source Routing**

- **Node-Table Routing**

- **Combinational Circuits**

- *To be discussed when we discuss router microarchitecture!*

Recap: O1TURN Routing Algorithm

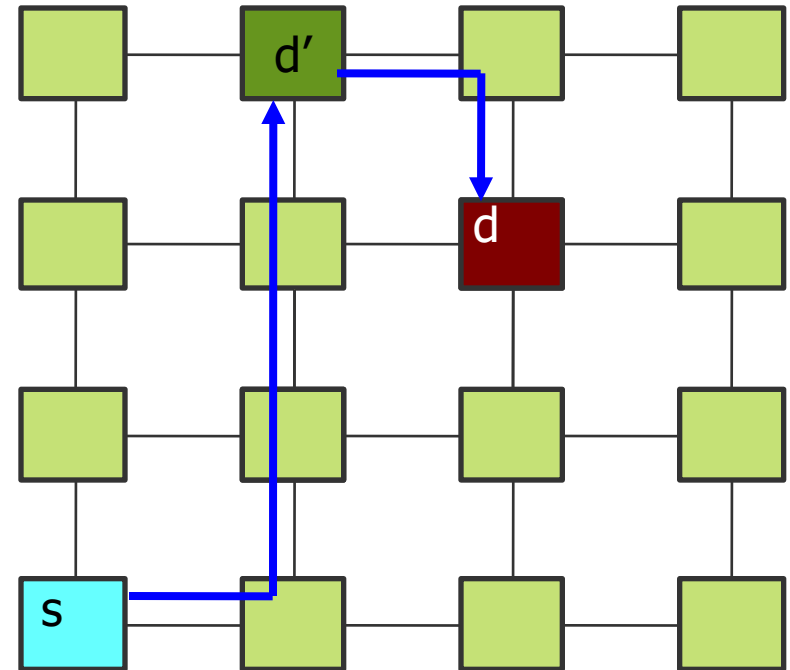


Randomly send over XY or YX

Minimal and Oblivious

Recap: Valiant's Routing Algorithm

- To route from s to d
 - Randomly choose intermediate node d'
 - Route* from s to d' (Phase I), and d' to d (Phase II)
- Pros
 - Randomizes any traffic pattern
 - All patterns appear uniform random
 - Balances network-load
 - Higher throughput
- Cons
 - Non-minimal
 - Higher latency and energy
 - Destroys locality




Non-Minimal and *Oblivious

****can also be Adaptive***

Deadlock


- A condition in which a set of **agents** wait indefinitely trying to acquire a set of **resources**



Note: holding buffer u == holding Channel 01 as no other packet can use channel 01 till buffer u becomes free

- Packet A holds buffer u (in 1) and wants buffer v (in 2)
- Packet B holds buffer v (in 2) and wants buffer w (in 3)
- Packet C holds buffer w (in 3) and wants buffer x (in 0)
- Packet D holds buffer x (in 0) and wants buffer u (in 1)

Classic Example: Dining Philosopher's Problem


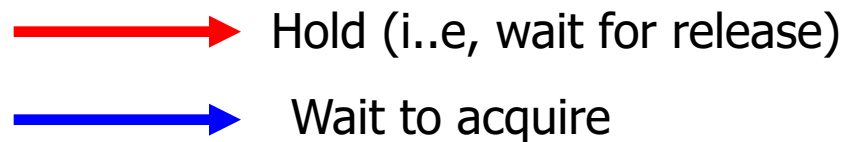


Agents:
Philosophers


Resources:
Forks

Resource Dependence

Resource A is **dependent** on resource B if it is possible for A to be *held-by* an agent X and it is also possible for X to *wait-for* B



Resource Dependence Graph



Deadlock Condition

- Agents hold and do not release a resource while waiting for access to another
- A cycle exists between waiting agents such that there exists a set of agents A_0, \dots, A_{n-1} , where agent A_i holds resource R_i , while waiting on resource $R_{(i+1 \bmod n)}$, for $i = 0, \dots, n-1$
- To avoid deadlock – resource dependence graph should not have any cycles

Dealing with Deadlocks

■ **Avoidance**

- Guarantee that the network will never deadlock
- Almost all modern networks use deadlock avoidance


■ **Recovery**

- Detect deadlock and correct

Deadlock Avoidance


- Eliminate cycles in Resource Dependency Graph
 - Resource Ordering
 - Enforce a partial/total order on the resources, and insist that an agent acquire the resources in ascending order
 - Deadlock avoided since a cycle must contain at least one agent holding a higher numbered resource waiting for a lower-numbered resource which is not allowed by the ordering allocation
 - Implementation
 - Restrict certain routes so that a higher numbered resource cannot wait for a lower numbered resource
 - Partition the buffers at each node such that they belong to different resource classes. A packet only any route can only acquire buffers in ascending order of resource class

Turn Model (Glass and Ni 1994) for Mesh




- Deadlocks may occur if the turns taken form a cycle
 - Removing some turns can make the routing algorithm deadlock free


Dimension Ordered Routing



XY Model




YX Model




O1Turn Deadlock!

Deadlock-free Oblivious/Adaptive Routing Algorithms


16



West-First Turn Model

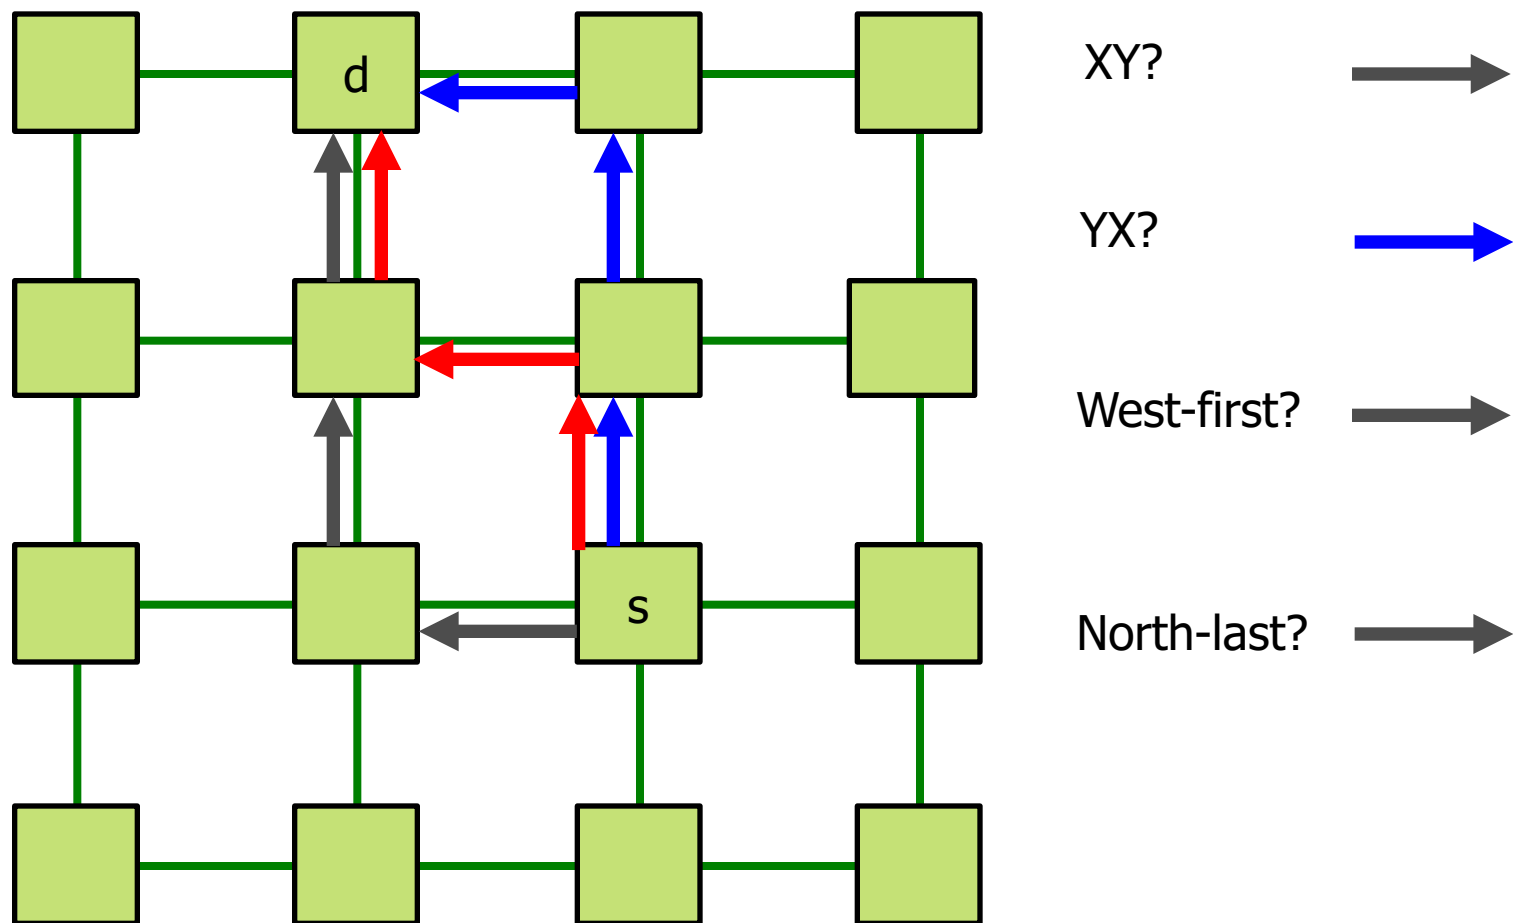


North-Last Turn Model

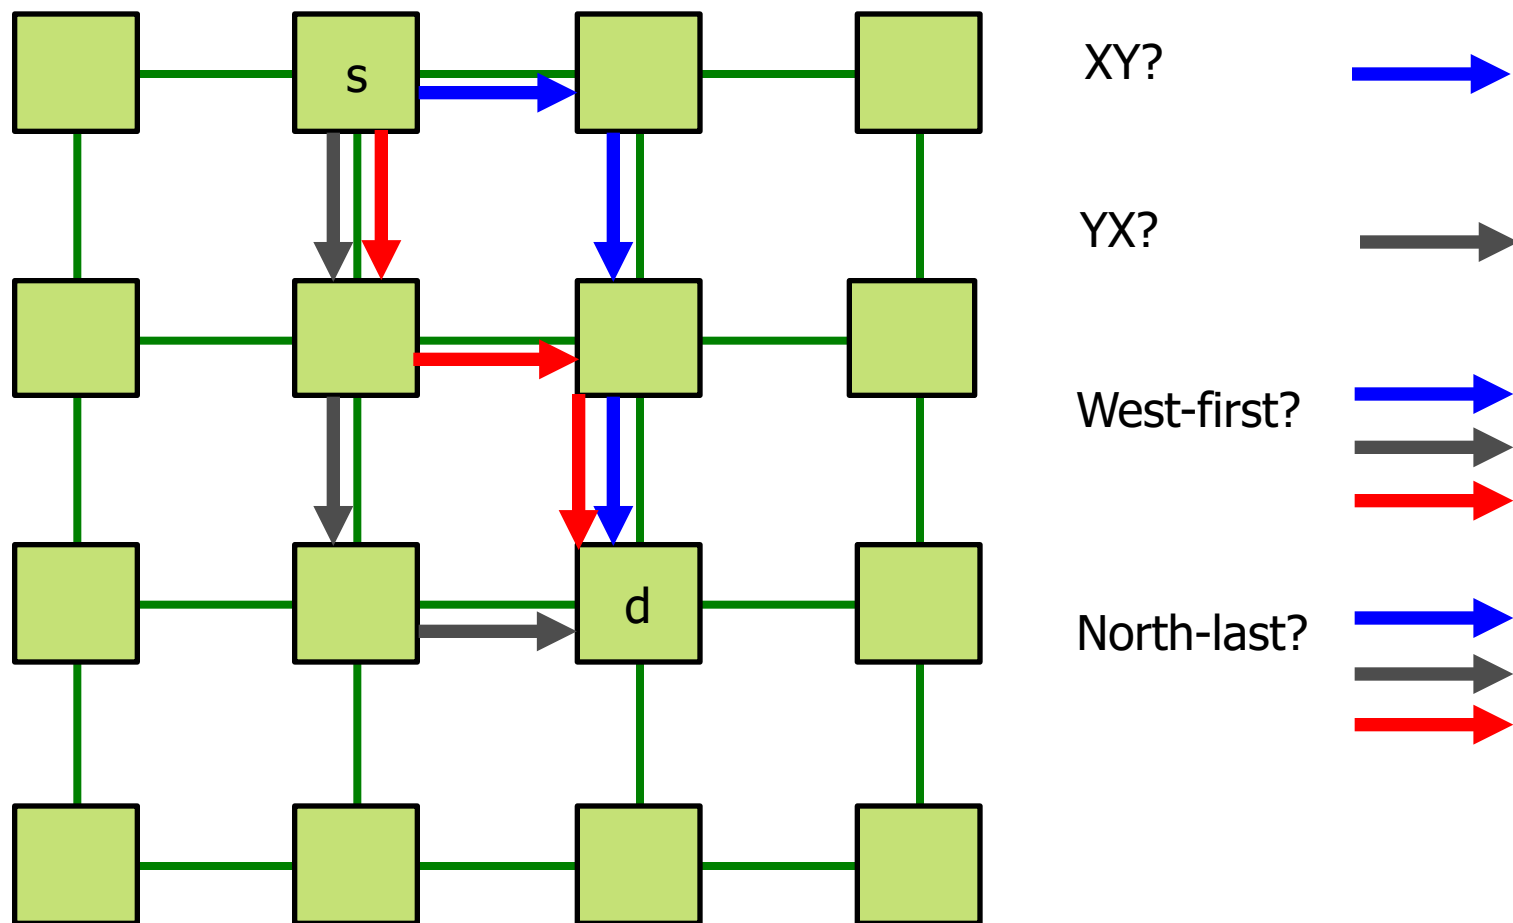


Negative-First Turn Model


Example 1




Example 2



Resource (channel) Ordering




XY Model




West-First Turn Model

Can we eliminate *any* 2 turns?



Six turn model



Deadlock!



Total Turn Models = 16

Deadlock Free = 12

Unique (non-symmetrical) = 3



Channel Dependency Graph (CDG)

- Vertices represent network links (channels)
- Edges represent turns
 - *180° turns not allowed, e.g., $AB \rightarrow BA$*



Cycles in the CDG


The channel dependency graph D derived from the network topology may contain many *cycles*



Flow routed through links AB, BE, EF
 Flow routed through links EF, FA, AB
 Deadlock!


Edges in CDG = Turns in Network
 → Disallow/Delete certain edges in CDG

Acyclic CDG




*This is the
West-first turn model!*

Cyclic CDG




↓ *Disable certain edges*




Acyclic CDG

CDG for arbitrary topology




Deadlock free?

No

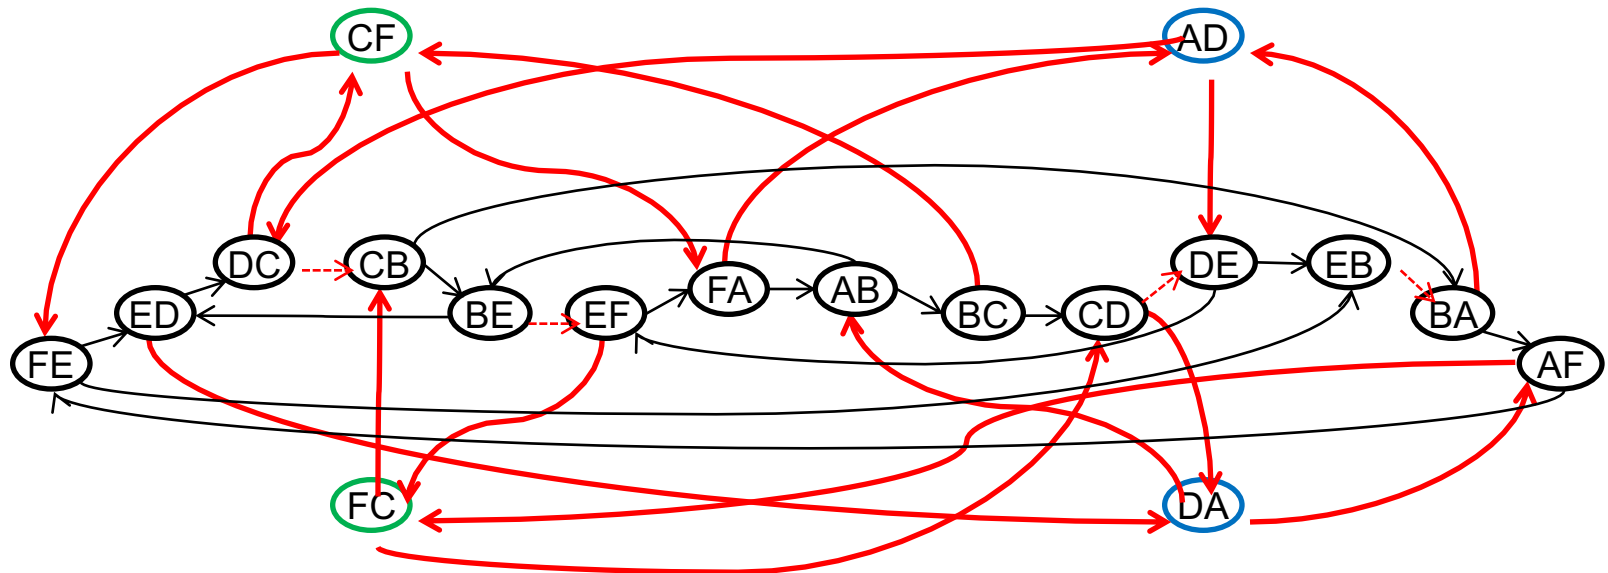


CDG for West-first
turn model


Deadlock-free Routing Algorithm



Suppose: Diagonal links should be traversed last (i.e., no edge from blue/green channel to black)



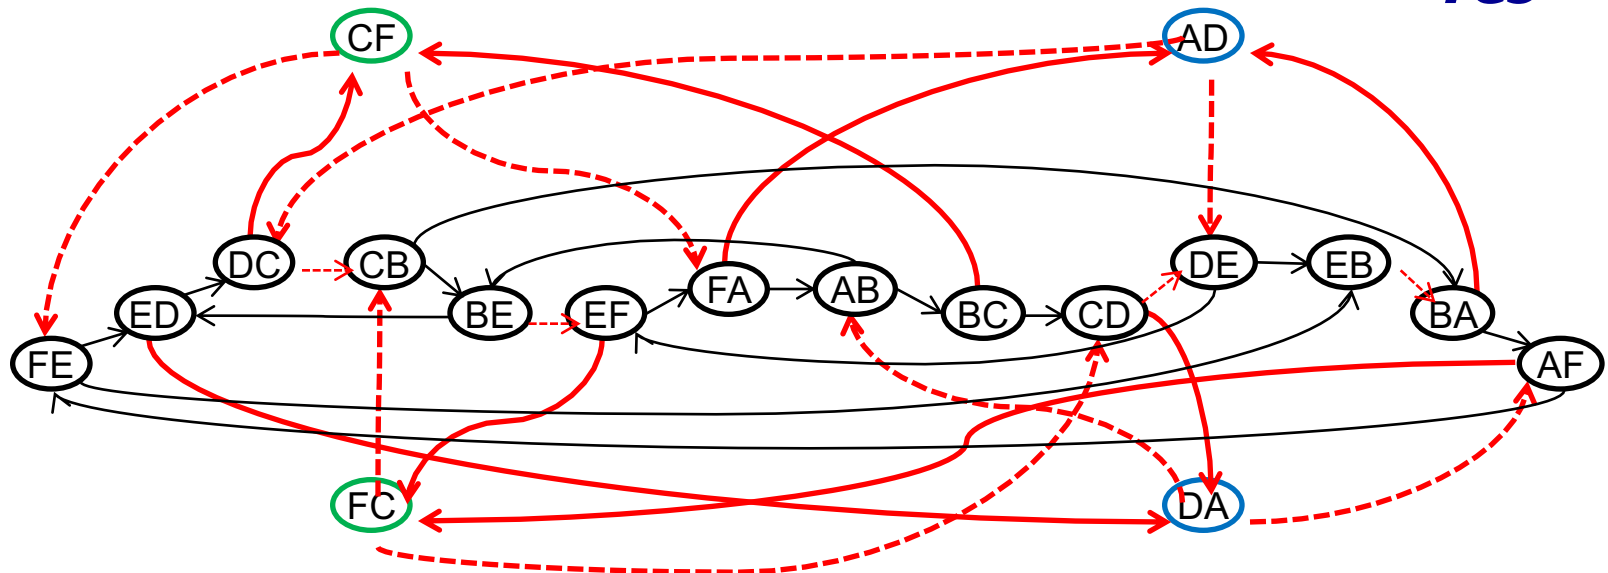
Deadlock-free Routing Algorithm




Suppose: Diagonal links should be traversed last (i.e., no edge from blue/green channel to black)

Deadlock free?


Yes



What about a Ring?




Acyclic CDG for a Ring




Route from E to C disabled
(E to D) and (D to C) allowed

Route from F to B disabled

Option 1




CDG



Problem? No route from E/F to B/C


Acyclic CDG for a Ring




Route from E to C disabled
(E to D) and (D to C) allowed

Route from E to A disabled

Option 2




CDG



Problem? No route from E to A/B/C


Acyclic CDG for a Ring




Route from E to C disabled
(E to D) and (D to C) allowed

Route from B to D disabled

Option 3




CDG



Acceptable CDG

Problem? E to C no longer minimal

Acyclic CDG for a Large Ring

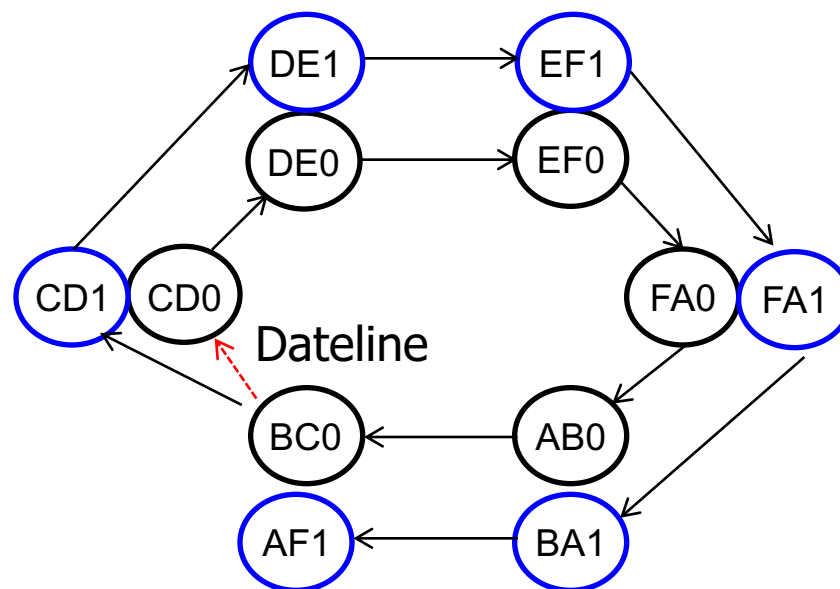
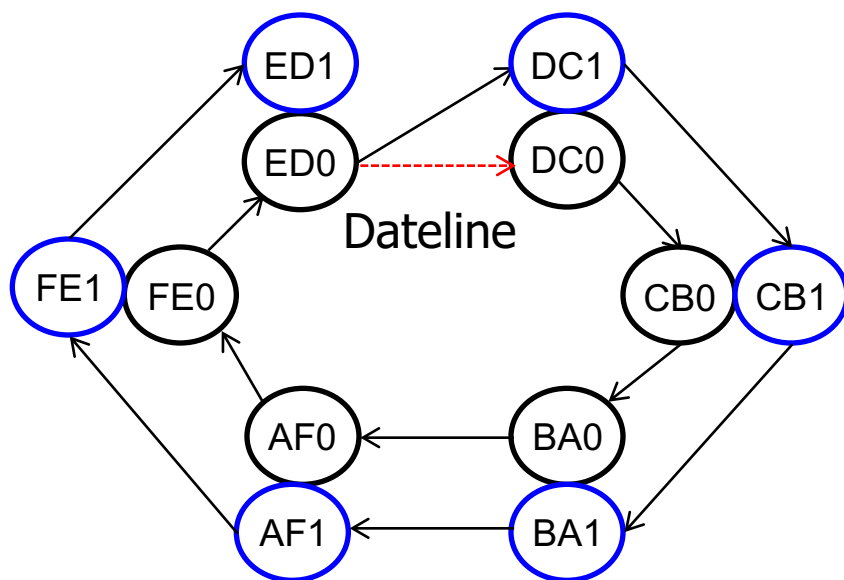
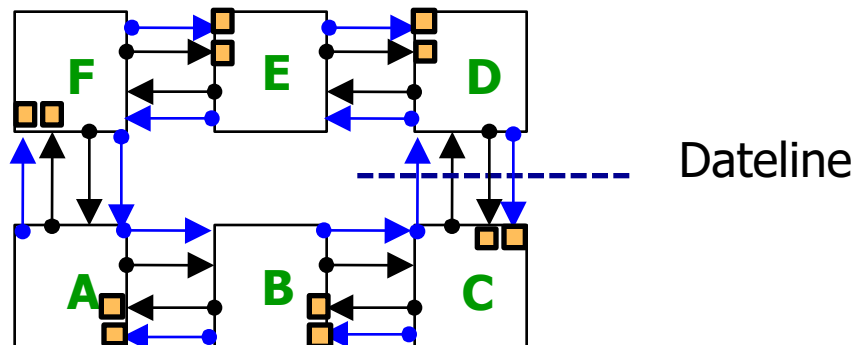


Problem?

G, H, I have to take non-minimal paths
to reach E!

D, C, B have to take non-minimal paths
to reach F

Suppose *two* channels



Need not be physical channels

Need at least 2 classes of buffers - called **"Virtual Channels"**

*Start in VC in Class0
After Dateline, jump to VC in Class1*

